

Importing the libraries

```
In [157... import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [158... cd
```

C:\Users\rusha

Reading the dataset

```
In [159... df = pd.read_csv(r"C:\Users\rusha\Downloads\data.csv")
```

```
In [160... df.head()
```

```
Out[160]:
```

	age	job	marital	education	default	balance	housing	loan	contact	day	month
0	38	technician	married	tertiary	no	127	yes	no	cellular	14	oct
1	41	housemaid	married	primary	no	365	no	no	cellular	8	aug
2	39	management	single	tertiary	no	2454	yes	no	cellular	4	may
3	49	blue-collar	married	primary	no	6215	yes	no	cellular	11	may
4	37	services	married	secondary	no	1694	yes	yes	cellular	29	jan

Checking the data


```
In [161... count=df['deposit'].value_counts()
```

```
In [162... count
```

```
Out[162]: no    4428
yes    3943
Name: deposit, dtype: int64
```

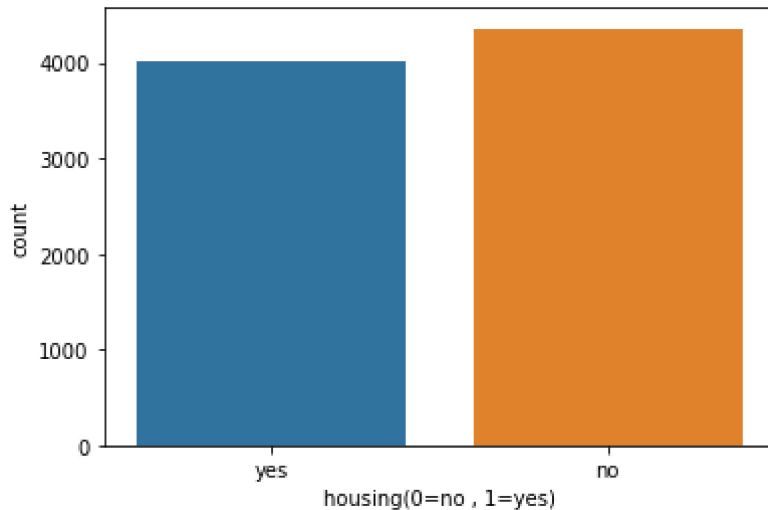
```
In [163... count.plot(kind="bar")
```

```
Out[163]: <AxesSubplot:>
```



```
In [164... sns.countplot(x = "housing", data=df)
plt.xlabel("housing(0=no , 1=yes)")
```

Out[164]: Text(0.5, 0, 'housing(0=no , 1=yes)')



```
In [165... count=df['housing'].value_counts()
count
```

Out[165]: no 4354
yes 4017
Name: housing, dtype: int64

We conclude from the above plot that

1. 4354 individuals did not deposit the money
2. 4017 individuals deposited the money

Applying Machine Learning Techniques

Logistic Regression

```
In [166... from sklearn import linear_model
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
```

```
In [167... le=preprocessing.LabelEncoder()
df=df.apply(le.fit_transform)
```

```
In [168... x= df.iloc[:, :-1].values
y= df.iloc[:, -1].values
```

```
In [169... x_train, x_test, y_train, y_test = train_test_split(x,y,test_size = 0.2,random_sta
```

```
In [170... lr=linear_model.LogisticRegression()
lr.fit(x_train,y_train)
```

```
D:\machinelearning\lib\site-packages\sklearn\linear_model\_logistic.py:814: Conver
genceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
```

```
Out[170]: LogisticRegression()
```

```
In [171... y_pred=lr.predict(x_test)
```

```
In [172... y_pred
```

```
Out[172]: array([0, 0, 1, ..., 0, 1, 1])
```

Confusion Matrix

```
In [173... from sklearn.metrics import confusion_matrix
cm=confusion_matrix(y_test,y_pred)
```

```
In [174... cm
```

```
Out[174]: array([[716, 183],
 [213, 563]], dtype=int64)
```

```
In [175... confusion_matrix(y_train,y_train)
```

```
Out[175]: array([[3529,  0],
 [  0, 3167]], dtype=int64)
```

```
In [176... from sklearn.metrics import precision_score , recall_score
```

```
In [177... precision_score(y_test,y_pred)
```

```
Out[177]: 0.7546916890080428
```

```
In [178... recall_score(y_test,y_pred)
```

```
Out[178]: 0.7255154639175257
```

```
In [179... from sklearn.metrics import f1_score
```

```
In [180... f1_score(y_test,y_pred)
```

```
Out[180]: 0.7398160315374506
```

```
In [181... accuracy=(716+563)/(716+183+213+563)
```

```
In [182... accuracy*100
```

```
Out[182]: 76.35820895522389
```

```
In [183... from sklearn.metrics import classification_report
```

```
In [184... print(classification_report(y_test,y_pred))
```

	resultofMLMarathon			
	precision	recall	f1-score	support
0	0.77	0.80	0.78	899
1	0.75	0.73	0.74	776
accuracy			0.76	1675
macro avg	0.76	0.76	0.76	1675
weighted avg	0.76	0.76	0.76	1675

Decision Tree Matrix

```
In [185... from sklearn.tree import DecisionTreeClassifier
dtc=DecisionTreeClassifier()
dtc.fit(x_train,y_train)
print("Decision Tree Test Accuracy {:.2f}%".format(dtc.score(x_test,y_test)*100))
```

Decision Tree Test Accuracy 75.34%

Random Forest Matrix

```
In [186... from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(n_estimators = 1000, random_state = 1)
rf.fit(x_train, y_train)
print("Random Forest Algorithm Accuracy Score : {:.2f}%".format(rf.score(x_test,y_
```

Random Forest Algorithm Accuracy Score : 83.52%